

```
import keras
keras.__version__

'2.12.0'

from keras.layers import Embedding

# The Embedding layer takes at least two arguments:
# the number of possible tokens, here 1000 (1 + maximum word index),
# and the dimensionality of the embeddings, here 64.
embedding_layer = Embedding(1000, 64)

from keras.datasets import imdb
from keras import preprocessing
from keras import utils as np_utils
#from tensorflow.keras.preprocessing.sequence import pad_sequences
#from keras.utils import pad_sequences
from keras.utils.data_utils import pad_sequences
# Number of words to consider as features
max_features = 10000
# Cut texts after this number of words
# (among top max_features most common words)
maxlen = 150

# Load the data as lists of integers.
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_features)

# This turns our lists of integers
# into a 2D integer tensor of shape `(samples, maxlen)`
x_train = keras.utils.pad_sequences(x_train, maxlen=maxlen)
x_test = keras.utils.pad_sequences(x_test, maxlen=maxlen)
```

```
print(len(x_train))
```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```
import numpy as np
from keras.models import Sequential
from keras.layers import Flatten, Dense

model = Sequential()
# We specify the maximum input length to our Embedding layer
# so we can later flatten the embedded inputs
model.add(Embedding(10000, 8, input_length=maxlen))
# After the Embedding layer,
# our activations have shape `(samples, maxlen, 8)`.

# We flatten the 3D tensor of embeddings
# into a 2D tensor of shape `(samples, maxlen * 8)`
model.add(Flatten())

# We add the classifier on top
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
model.summary()

history = model.fit(x_train, y_train,
                    epochs=10,
                    batch_size=32,
                    validation_split=0.2)
```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
embedding_11 (Embedding)	(None, 150, 8)	80000
flatten_8 (Flatten)	(None, 1200)	0
dense_10 (Dense)	(None, 1)	1201

=====
 Total params: 81,201
 Trainable params: 81,201
 Non-trainable params: 0

```

Epoch 1/10
625/625 [=====] - 27s 41ms/step - loss: 0.5914 - acc: 0.7071 - val_loss: 0.4237 - val_acc: 0.825
Epoch 2/10
625/625 [=====] - 7s 11ms/step - loss: 0.3310 - acc: 0.8651 - val_loss: 0.3211 - val_acc: 0.8651
Epoch 3/10
625/625 [=====] - 4s 7ms/step - loss: 0.2582 - acc: 0.8963 - val_loss: 0.3031 - val_acc: 0.8730
Epoch 4/10
625/625 [=====] - 3s 5ms/step - loss: 0.2229 - acc: 0.9134 - val_loss: 0.3024 - val_acc: 0.8712
Epoch 5/10
625/625 [=====] - 4s 7ms/step - loss: 0.1979 - acc: 0.9254 - val_loss: 0.3042 - val_acc: 0.8722
Epoch 6/10
625/625 [=====] - 3s 5ms/step - loss: 0.1782 - acc: 0.9338 - val_loss: 0.3102 - val_acc: 0.8694
Epoch 7/10
625/625 [=====] - 3s 5ms/step - loss: 0.1599 - acc: 0.9426 - val_loss: 0.3176 - val_acc: 0.8696
Epoch 8/10
625/625 [=====] - 3s 4ms/step - loss: 0.1424 - acc: 0.9495 - val_loss: 0.3258 - val_acc: 0.8694
Epoch 9/10
625/625 [=====] - 3s 4ms/step - loss: 0.1256 - acc: 0.9581 - val_loss: 0.3365 - val_acc: 0.8656
Epoch 10/10
625/625 [=====] - 4s 6ms/step - loss: 0.1099 - acc: 0.9640 - val_loss: 0.3530 - val_acc: 0.8630

```

```

from google.colab import drive
drive.mount('/content/gdrive')

```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force_remount=True)

```

import os
import shutil
mkdir ..data
!wget -O ../data/aclImdb_v1.tar.gz http://ai.stanford.edu/~amaas/data/sentiment/aclImdb_v1.tar.gz
!tar -zxvf ../data/aclImdb_v1.tar.gz -C ../data
imdb_dir = '/content/gdrive/MyDrive/ML Assignment 3/aclImdb'
train_dir = '/content/gdrive/MyDrive/ML Assignment 3/aclImdb/train'
#train_dir = os.path.join(imdb_dir, 'train')

```

```

labels = []
texts = []

```

```

for label_type in ['neg', 'pos']:
    dir_name = os.path.join(train_dir, label_type)

```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```

f = open(os.path.join(dir_name, fname))
texts.append(f.read())
f.close()
if label_type == 'neg':
    labels.append(0)
else:
    labels.append(1)

```

```

mkdir: cannot create directory '../data': File exists
--2023-04-16 19:11:07-- http://ai.stanford.edu/~amaas/data/sentiment/aclImdb_v1.tar.gz
Resolving ai.stanford.edu (ai.stanford.edu)... 171.64.68.10
Connecting to ai.stanford.edu (ai.stanford.edu)|171.64.68.10|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 84125825 (80M) [application/x-gzip]
Saving to: '../data/aclImdb_v1.tar.gz'

```

```
../data/aclImdb_v1. 100%[=====] 80.23M 67.1MB/s in 1.2s
```

```
2023-04-16 19:11:08 (67.1 MB/s) - '../data/aclImdb_v1.tar.gz' saved [84125825/84125825]
```

```

from keras.preprocessing.text import Tokenizer
from keras.utils.data_utils import pad_sequences
#from keras.preprocessing.sequence import pad_sequences
import numpy as np

```

```

maxlen = 150 # We will cut reviews after 150 words
training_samples = 100 # We will be training on 100 samples
validation_samples = 10000 # We will be validating on 10000 samples
max_words = 10000 # We will only consider the top 10,000 words in the dataset

```

```

tokenizer = Tokenizer(num_words=max_words)
tokenizer.fit_on_texts(texts)
sequences = tokenizer.texts_to_sequences(texts)

```

```

word_index = tokenizer.word_index
print('Found %s unique tokens.' % len(word_index))

```

```
data = pad_sequences(sequences, maxlen=maxlen)
```

```

labels = np.asarray(labels)
print('Shape of data tensor:', data.shape)
print('Shape of label tensor:', labels.shape)

# Split the data into a training set and a validation set
# But first, shuffle the data, since we started from data
# where sample are ordered (all negative first, then all positive).
indices = np.arange(data.shape[0])
np.random.shuffle(indices)
data = data[indices]
labels = labels[indices]

x_train = data[:training_samples]
y_train = labels[:training_samples]
x_val = data[training_samples: training_samples + validation_samples]
y_val = labels[training_samples: training_samples + validation_samples]

Found 88582 unique tokens.
Shape of data tensor: (25000, 150)
Shape of label tensor: (25000,)

```

```

import os
import numpy as np
from keras.preprocessing.text import Tokenizer
from keras.utils.data_utils import pad_sequences
glove_dir = '/content/gdrive/MyDrive/ML Assignment 3/glove6B'

embeddings_index = {}
f = open(os.path.join(glove_dir, 'glove.6B.100d.txt'))
for line in f:
    values = line.split()
    word = values[0]
    coefs = np.asarray(values[1:], dtype='float32')
    embeddings_index[word] = coefs
f.close()

print('Found %s word vectors.' % len(embeddings_index))

Found 400001 word vectors.

```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```

embedding_dim = 100
max_words = 10000
tokenizer = Tokenizer(num_words=max_words)

word_index = tokenizer.word_index
print('Found %s unique tokens.' % len(word_index))
embedding_matrix = np.zeros((max_words, embedding_dim))
for word, i in word_index.items():
    embedding_vector = embeddings_index.get(word)
    if i < max_words:
        if embedding_vector is not None:
            # Words not found in embedding index will be all-zeros.
            embedding_matrix[i] = embedding_vector

Found 0 unique tokens.

```

```

from keras.models import Sequential
from keras.layers import Embedding, Flatten, Dense
maxlen = 150
model = Sequential()
model.add(Embedding(max_words, embedding_dim, input_length=maxlen))
model.add(Flatten())
model.add(Dense(32, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.summary()

```

Model: "sequential_9"

Layer (type)	Output Shape	Param #
embedding_12 (Embedding)	(None, 150, 100)	1000000
flatten_9 (Flatten)	(None, 15000)	0
dense_11 (Dense)	(None, 32)	480032
dense_12 (Dense)	(None, 1)	33
Total params: 1,480,065		

```
Trainable params: 1,480,065
Non-trainable params: 0
```

```
model.layers[0].set_weights([embedding_matrix])
model.layers[0].trainable = False

model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['acc'])
history = model.fit(x_train, y_train,
                   epochs=10,
                   batch_size=32,
                   validation_data=(x_val, y_val))
model.save_weights('pre_trained_glove_model.h5')
```

```
Epoch 1/10
4/4 [=====] - 2s 353ms/step - loss: 0.6929 - acc: 0.5900 - val_loss: 0.6931 - val_acc: 0.5060
Epoch 2/10
4/4 [=====] - 1s 196ms/step - loss: 0.6923 - acc: 0.5900 - val_loss: 0.6931 - val_acc: 0.5060
Epoch 3/10
4/4 [=====] - 1s 218ms/step - loss: 0.6921 - acc: 0.5900 - val_loss: 0.6931 - val_acc: 0.5060
Epoch 4/10
4/4 [=====] - 1s 218ms/step - loss: 0.6921 - acc: 0.5900 - val_loss: 0.6931 - val_acc: 0.5060
Epoch 5/10
4/4 [=====] - 1s 290ms/step - loss: 0.6919 - acc: 0.5900 - val_loss: 0.6931 - val_acc: 0.5060
Epoch 6/10
4/4 [=====] - 1s 289ms/step - loss: 0.6918 - acc: 0.5900 - val_loss: 0.6931 - val_acc: 0.5060
Epoch 7/10
4/4 [=====] - 1s 290ms/step - loss: 0.6916 - acc: 0.5900 - val_loss: 0.6931 - val_acc: 0.5060
Epoch 8/10
4/4 [=====] - 1s 437ms/step - loss: 0.6914 - acc: 0.5900 - val_loss: 0.6931 - val_acc: 0.5060
Epoch 9/10
4/4 [=====] - 1s 219ms/step - loss: 0.6912 - acc: 0.5900 - val_loss: 0.6931 - val_acc: 0.5060
Epoch 10/10
4/4 [=====] - 1s 219ms/step - loss: 0.6910 - acc: 0.5900 - val_loss: 0.6931 - val_acc: 0.5060
```

```
import matplotlib.pyplot as plt
```

```
acc = history.history['acc']
```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```
val_loss = history.history['val_loss']
```


```
epochs = range(1, len(acc) + 1)
```

```
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
```

```
plt.figure()
```

```
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
```

```
plt.show()
```



```
test_dir = os.path.join(imdb_dir, 'test')
```

```
labels = []
texts = []
```

```
for label_type in ['neg', 'pos']:
    dir_name = os.path.join(test_dir, label_type)
    for fname in sorted(os.listdir(dir_name)):
        if fname[-4:] == '.txt':
            f = open(os.path.join(dir_name, fname))
            texts.append(f.read())
            f.close()
            if label_type == 'neg':
                labels.append(0)
            else:
                labels.append(1)
```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```
x_test = pad_sequences(sequences, maxlen=maxlen)
y_test = np.asarray(labels)
```

```
model.load_weights('pre_trained_glove_model.h5')
model.evaluate(x_test, y_test)
```

```
785/785 [=====] - 2s 3ms/step - loss: 0.6933 - acc: 0.4979
[0.6932905316352844, 0.49794843792915344]
```

Hypertuning Embedding Layer 1 - 1000 Samples

```
import keras
keras.__version__

'2.12.0'
```

```
from keras.layers import Embedding
```

```
# The Embedding layer takes at least two arguments:
# the number of possible tokens, here 1000 (1 + maximum word index),
# and the dimensionality of the embeddings, here 64.
embedding_layer = Embedding(1000, 64)
```

```
from keras.datasets import imdb
from keras import preprocessing
from keras.utils.data_utils import pad_sequences
```

```
# Number of words to consider as features
max_features = 10000
# Cut texts after this number of words
# (among top max_features most common words)
maxlen = 150
```

```
# Load the data as lists of integers.
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_features)
```

```
x_train = x_train[:1000]
y_train = y_train[:1000]

# This turns our lists of integers
# into a 2D integer tensor of shape `(samples, maxlen)`
x_train = keras.utils.pad_sequences(x_train, maxlen=maxlen)
x_test = keras.utils.pad_sequences(x_test, maxlen=maxlen)

print(len(x_train))

1000

from keras.models import Sequential
from keras.layers import Flatten, Dense

model = Sequential()
# We specify the maximum input length to our Embedding layer
# so we can later flatten the embedded inputs
model.add(Embedding(10000, 8, input_length=maxlen))
# After the Embedding layer,
# our activations have shape `(samples, maxlen, 8)`.

# We flatten the 3D tensor of embeddings
# into a 2D tensor of shape `(samples, maxlen * 8)`
model.add(Flatten())

# We add the classifier on top
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
model.summary()

history = model.fit(x_train, y_train,
                    epochs=10,
                    batch_size=32,
```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

Layer (type)	Output Shape	Param #
embedding_14 (Embedding)	(None, 150, 8)	80000
flatten_10 (Flatten)	(None, 1200)	0
dense_13 (Dense)	(None, 1)	1201

=====
Total params: 81,201
Trainable params: 81,201
Non-trainable params: 0

```
Epoch 1/10
25/25 [=====] - 3s 110ms/step - loss: 0.6939 - acc: 0.5000 - val_loss: 0.6921 - val_acc: 0.5200
Epoch 2/10
25/25 [=====] - 3s 127ms/step - loss: 0.6771 - acc: 0.7563 - val_loss: 0.6904 - val_acc: 0.5250
Epoch 3/10
25/25 [=====] - 3s 105ms/step - loss: 0.6607 - acc: 0.8712 - val_loss: 0.6887 - val_acc: 0.5250
Epoch 4/10
25/25 [=====] - 2s 73ms/step - loss: 0.6401 - acc: 0.9250 - val_loss: 0.6865 - val_acc: 0.5100
Epoch 5/10
25/25 [=====] - 2s 67ms/step - loss: 0.6147 - acc: 0.9400 - val_loss: 0.6837 - val_acc: 0.5350
Epoch 6/10
25/25 [=====] - 1s 53ms/step - loss: 0.5843 - acc: 0.9500 - val_loss: 0.6804 - val_acc: 0.5500
Epoch 7/10
25/25 [=====] - 1s 40ms/step - loss: 0.5490 - acc: 0.9613 - val_loss: 0.6763 - val_acc: 0.5600
Epoch 8/10
25/25 [=====] - 1s 41ms/step - loss: 0.5095 - acc: 0.9712 - val_loss: 0.6717 - val_acc: 0.5800
Epoch 9/10
25/25 [=====] - 1s 31ms/step - loss: 0.4666 - acc: 0.9725 - val_loss: 0.6665 - val_acc: 0.5900
Epoch 10/10
25/25 [=====] - 1s 44ms/step - loss: 0.4218 - acc: 0.9812 - val_loss: 0.6611 - val_acc: 0.6050
```

800 Samples

```

from keras.datasets import imdb
from keras import preprocessing
from keras.utils.data_utils import pad_sequences

# Number of words to consider as features
max_features = 10000
# Cut texts after this number of words
# (among top max_features most common words)
maxlen = 150

# Load the data as lists of integers.
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_features)

x_train = x_train[:800]
y_train = y_train[:800]

# This turns our lists of integers
# into a 2D integer tensor of shape `(samples, maxlen)`
x_train = keras.utils.pad_sequences(x_train, maxlen=maxlen)
x_test = keras.utils.pad_sequences(x_test, maxlen=maxlen)

print(len(x_train))

800

from keras.models import Sequential
from keras.layers import Flatten, Dense

model = Sequential()
# We specify the maximum input length to our Embedding layer
# so we can later flatten the embedded inputs
model.add(Embedding(10000, 8, input_length=maxlen))
# After the Embedding layer,
# our activations have shape `(samples, maxlen, 8)`.

# We flatten the 3D tensor of embeddings
# into a 2D tensor of shape `(samples, maxlen * 8)`
model.add(Flatten())

Automatic saving failed. This file was updated remotely or in another tab. Show diff

model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
model.summary()

history = model.fit(x_train, y_train,
                    epochs=10,
                    batch_size=32,
                    validation_split=0.2)

```

Model: "sequential_11"

Layer (type)	Output Shape	Param #
=====		
embedding_15 (Embedding)	(None, 150, 8)	80000
flatten_11 (Flatten)	(None, 1200)	0
dense_14 (Dense)	(None, 1)	1201
=====		
Total params: 81,201		
Trainable params: 81,201		
Non-trainable params: 0		

```

Epoch 1/10
20/20 [=====] - 3s 107ms/step - loss: 0.6942 - acc: 0.4828 - val_loss: 0.6905 - val_acc: 0.5250
Epoch 2/10
20/20 [=====] - 2s 99ms/step - loss: 0.6767 - acc: 0.7188 - val_loss: 0.6898 - val_acc: 0.5188
Epoch 3/10
20/20 [=====] - 2s 84ms/step - loss: 0.6602 - acc: 0.8359 - val_loss: 0.6887 - val_acc: 0.5250
Epoch 4/10
20/20 [=====] - 2s 123ms/step - loss: 0.6405 - acc: 0.8984 - val_loss: 0.6872 - val_acc: 0.5500
Epoch 5/10
20/20 [=====] - 2s 96ms/step - loss: 0.6167 - acc: 0.9141 - val_loss: 0.6847 - val_acc: 0.5437
Epoch 6/10
20/20 [=====] - 2s 77ms/step - loss: 0.5886 - acc: 0.9531 - val_loss: 0.6821 - val_acc: 0.5562
Epoch 7/10
20/20 [=====] - 1s 54ms/step - loss: 0.5563 - acc: 0.9641 - val_loss: 0.6788 - val_acc: 0.5875
Epoch 8/10
20/20 [=====] - 1s 44ms/step - loss: 0.5201 - acc: 0.9656 - val_loss: 0.6746 - val_acc: 0.5938
Epoch 9/10
20/20 [=====] - 1s 33ms/step - loss: 0.4809 - acc: 0.9672 - val_loss: 0.6698 - val_acc: 0.6125

```

```
Epoch 10/10
20/20 [=====] - 1s 21ms/step - loss: 0.4394 - acc: 0.9750 - val_loss: 0.6651 - val_acc: 0.6125
```

700 Samples

```
from keras.datasets import imdb
from keras import preprocessing

# Number of words to consider as features
max_features = 10000
# Cut texts after this number of words
# (among top max_features most common words)
maxlen = 150

# Load the data as lists of integers.
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_features)

x_train = x_train[:700]
y_train = y_train[:700]

# This turns our lists of integers
# into a 2D integer tensor of shape `(samples, maxlen)`
x_train = keras.utils.pad_sequences(x_train, maxlen=maxlen)
x_test = keras.utils.pad_sequences(x_test, maxlen=maxlen)

from keras.models import Sequential
from keras.layers import Flatten, Dense

model = Sequential()
# We specify the maximum input length to our Embedding layer
# so we can later flatten the embedded inputs
model.add(Embedding(10000, 8, input_length=maxlen))
# After the Embedding layer,
# our activations have shape `(samples, maxlen, 8)`.

# We flatten the 3D tensor of embeddings
# into a 2D tensor of shape `(samples, maxlen * 8)`
model.add(Flatten())

Automatic saving failed. This file was updated remotely or in another tab. Show diff

model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
model.summary()

history = model.fit(x_train, y_train,
                    epochs=10,
                    batch_size=32,
                    validation_split=0.2)
```

Model: "sequential_12"

Layer (type)	Output Shape	Param #
=====		
embedding_16 (Embedding)	(None, 150, 8)	80000
flatten_12 (Flatten)	(None, 1200)	0
dense_15 (Dense)	(None, 1)	1201
=====		
Total params: 81,201		
Trainable params: 81,201		
Non-trainable params: 0		

```
Epoch 1/10
18/18 [=====] - 2s 102ms/step - loss: 0.6926 - acc: 0.5089 - val_loss: 0.6962 - val_acc: 0.5071
Epoch 2/10
18/18 [=====] - 2s 111ms/step - loss: 0.6747 - acc: 0.7482 - val_loss: 0.6970 - val_acc: 0.5071
Epoch 3/10
18/18 [=====] - 2s 87ms/step - loss: 0.6588 - acc: 0.8607 - val_loss: 0.6974 - val_acc: 0.5214
Epoch 4/10
18/18 [=====] - 2s 85ms/step - loss: 0.6399 - acc: 0.9018 - val_loss: 0.6983 - val_acc: 0.4857
Epoch 5/10
18/18 [=====] - 2s 114ms/step - loss: 0.6179 - acc: 0.9339 - val_loss: 0.6994 - val_acc: 0.4857
Epoch 6/10
18/18 [=====] - 2s 129ms/step - loss: 0.5920 - acc: 0.9357 - val_loss: 0.7002 - val_acc: 0.4929
Epoch 7/10
18/18 [=====] - 1s 64ms/step - loss: 0.5626 - acc: 0.9482 - val_loss: 0.7005 - val_acc: 0.5214
Epoch 8/10
18/18 [=====] - 1s 57ms/step - loss: 0.5296 - acc: 0.9571 - val_loss: 0.7011 - val_acc: 0.5214
Epoch 9/10
18/18 [=====] - 1s 67ms/step - loss: 0.4943 - acc: 0.9571 - val_loss: 0.7017 - val_acc: 0.5214
```



```
Epoch 10/10
18/18 [=====] - 1s 51ms/step - loss: 0.4572 - acc: 0.9607 - val_loss: 0.7023 - val_acc: 0.5143
```

900 Samples

```
from keras.datasets import imdb
from keras import preprocessing

# Number of words to consider as features
max_features = 10000
# Cut texts after this number of words
# (among top max_features most common words)
maxlen = 150

# Load the data as lists of integers.
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_features)

x_train = x_train[700:1600]
y_train = y_train[700:1600]

# This turns our lists of integers
# into a 2D integer tensor of shape `(samples, maxlen)`
x_train = keras.utils.pad_sequences(x_train, maxlen=maxlen)
x_test = keras.utils.pad_sequences(x_test, maxlen=maxlen)

from keras.models import Sequential
from keras.layers import Flatten, Dense

model = Sequential()
# We specify the maximum input length to our Embedding layer
# so we can later flatten the embedded inputs
model.add(Embedding(10000, 8, input_length=maxlen))
# After the Embedding layer,
# our activations have shape `(samples, maxlen, 8)`.

# We flatten the 3D tensor of embeddings
# into a 2D tensor of shape `(samples, maxlen * 8)`
model.add(Flatten())

Automatic saving failed. This file was updated remotely or in another tab. Show diff

model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
model.summary()

history = model.fit(x_train, y_train,
                    epochs=10,
                    batch_size=32,
                    validation_split=0.2)
```

Model: "sequential_13"

Layer (type)	Output Shape	Param #
embedding_17 (Embedding)	(None, 150, 8)	80000
flatten_13 (Flatten)	(None, 1200)	0
dense_16 (Dense)	(None, 1)	1201
Total params: 81,201		
Trainable params: 81,201		
Non-trainable params: 0		

```
Epoch 1/10
23/23 [=====] - 3s 105ms/step - loss: 0.6948 - acc: 0.4750 - val_loss: 0.6949 - val_acc: 0.4556
Epoch 2/10
23/23 [=====] - 2s 98ms/step - loss: 0.6780 - acc: 0.7486 - val_loss: 0.6957 - val_acc: 0.4778
Epoch 3/10
23/23 [=====] - 2s 76ms/step - loss: 0.6620 - acc: 0.8833 - val_loss: 0.6969 - val_acc: 0.4556
Epoch 4/10
23/23 [=====] - 1s 59ms/step - loss: 0.6423 - acc: 0.9486 - val_loss: 0.6978 - val_acc: 0.4611
Epoch 5/10
23/23 [=====] - 1s 66ms/step - loss: 0.6181 - acc: 0.9792 - val_loss: 0.6987 - val_acc: 0.4500
Epoch 6/10
23/23 [=====] - 3s 125ms/step - loss: 0.5889 - acc: 0.9903 - val_loss: 0.6997 - val_acc: 0.4667
Epoch 7/10
23/23 [=====] - 1s 51ms/step - loss: 0.5553 - acc: 0.9972 - val_loss: 0.7013 - val_acc: 0.4667
Epoch 8/10
23/23 [=====] - 1s 64ms/step - loss: 0.5175 - acc: 0.9972 - val_loss: 0.7030 - val_acc: 0.4611
Epoch 9/10
23/23 [=====] - 2s 63ms/step - loss: 0.4765 - acc: 0.9972 - val_loss: 0.7041 - val_acc: 0.4556
```

Epoch 10/10

23/23 [=====] - 1s 62ms/step - loss: 0.4335 - acc: 0.9972 - val_loss: 0.7056 - val_acc: 0.4500



After using both an embedding layer and a pretrained layer, I determined that the plain embedding layer was the best approach (it was the simplest method, but it produced the greatest accuracy of 50.6% for the pretrained layer). After running some hypertuned embedding layers, I was able to determine that the regular embedding layer does better than the pretrained layer once the sample size reaches 800 samples. Anything below 800 samples, it would be best to use a pretrained word embedding layer. I ran the test originally with 1000 samples and pretrained layer, I was able to determine that the only way it was the simplest method, but the Accuracy incrementally increased and decreased the number of samples until I reached an accuracy slightly above the pretrained's accuracy of 45.00%. The pretrained technique is best suited for small sample sizes, but once the sample size reaches 900, the best method to use is the regular embedding layer.

After using both an embedding layer and a pretrained layer, I was able to determine that the plain embedding layer was the best approach. Not only was it the simplest method, but it produced the greatest accuracy (86.3% instead of 50.6% for the pretrained layer). After running some hypertuned embedding layers, I was able to determine that the regular embedding layer does better than the pretrained layer once the sample size reaches 800 samples. Anything below 800 samples, it would be best to use a pretrained word embedding layer. I ran the test originally with 1000 samples and pretrained layer, I was able to determine that the only way it was the simplest method, but it produced the Accuracy incrementally increased and decreased the number of samples until I reached an accuracy slightly above the pretrained's accuracy of 45.00%. The pretrained technique is best suited for small sample sizes, but once the sample size reaches 900, the best method to use is the regular embedding layer.

Automatic saving failed. This file was updated remotely or in another tab.

[Show diff](#)